



Most of the material has been taken from
Jurafsky & J. H. Martin: **Speech and Language
Processing**, Pearson International 3ed draft
<https://web.stanford.edu/~jurafsky/slp3/>

Jay Alammar, Maarten Grootendorst: Hands-
On Large Language Models Published by
O'Reilly Media, Inc.

Text Analysis -Transformers

Francesco Guerra
francesco.guerra@unimore.it

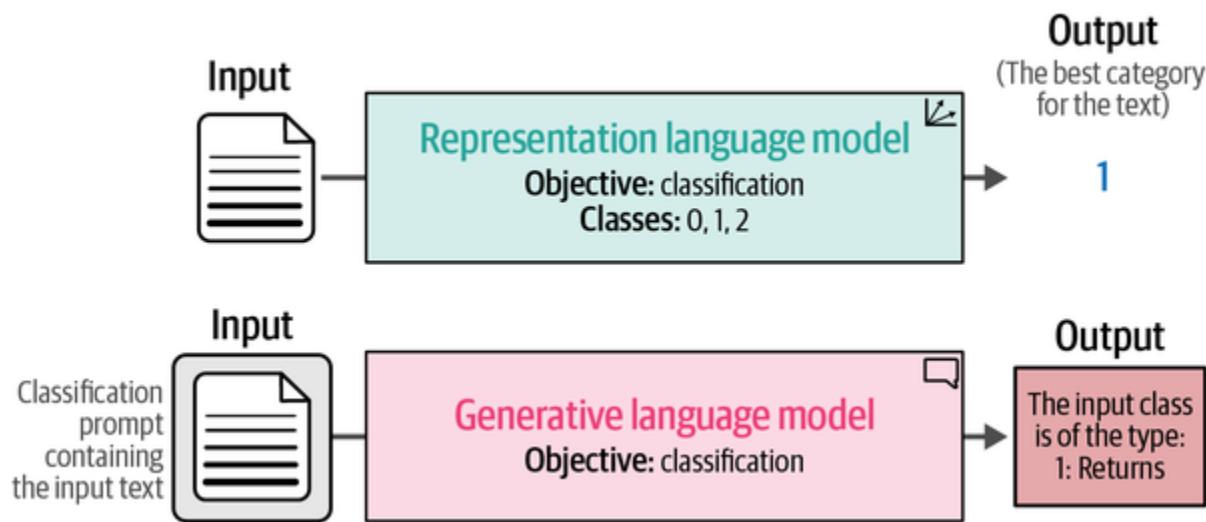
Big Data Analysis

Text Classification



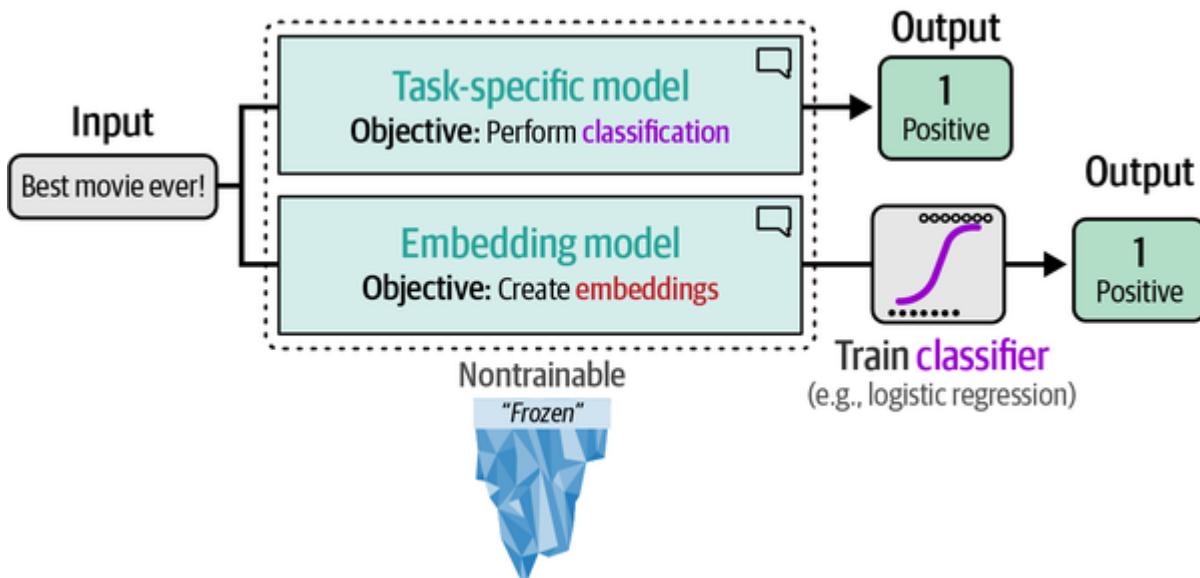
Text Classification

- ▶ The goal of the task is to train a model to assign a label or class to some input text.
 - ▶ Classifying text is used across for sentiment analysis, intent detection, extracting entities, detecting language ...
- ▶ Two base techniques:



Text Classification with Representation Models

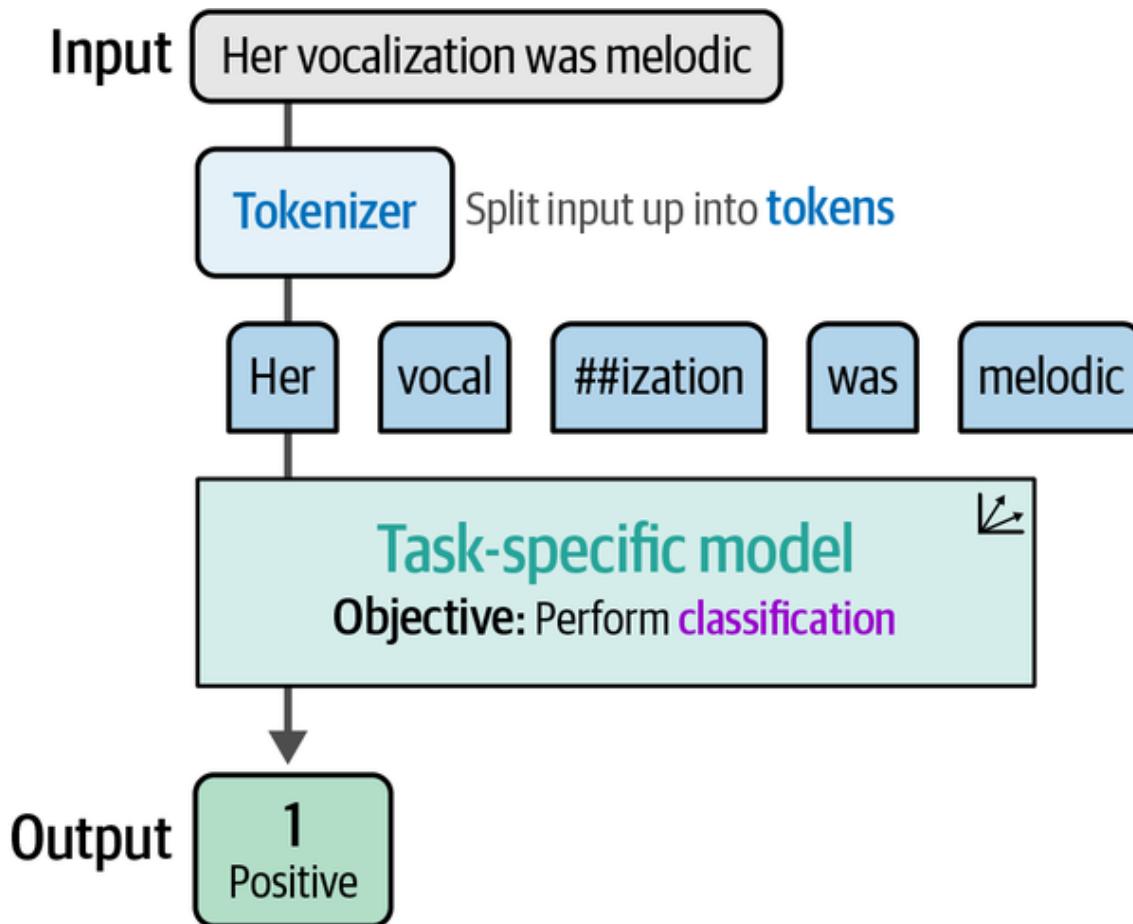
- ▶ Two techniques:
 - ▶ using a task-specific model or an embedding model



Model Selection

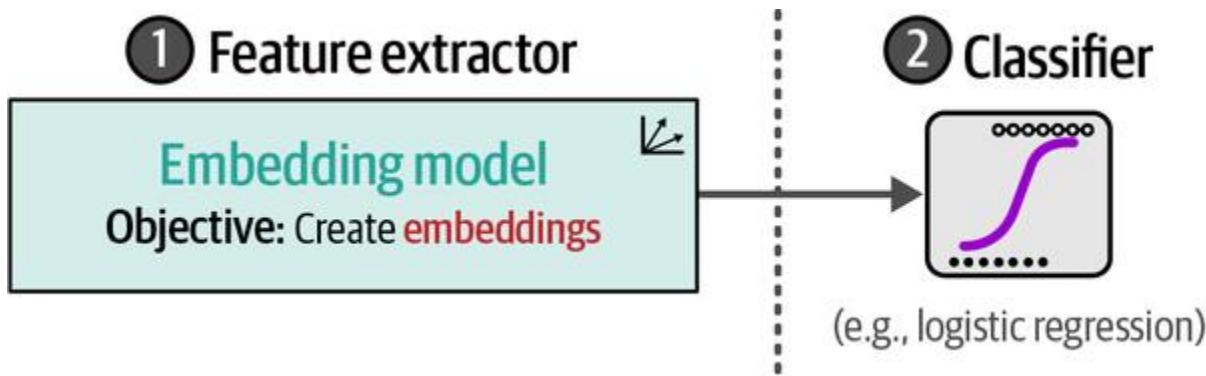
- ▶ Choosing the right models is not as straightforward
 - ▶ over 60,000 models on the Hugging Face Hub for text classification and more than 8,000 models that generate embeddings in 2025.
 - ▶ it's crucial to select a model that fits your use case and consider its language compatibility, the underlying architecture, size, and performance.
- ▶ BERT (or one of its variations) are a popular choice for creating task-specific and embedding models. <https://huggingface.co/models>
 - ▶ BERT base model (uncased)
 - ▶ RoBERTa base model
 - ▶ DistilBERT base model (uncased)
 - ▶ DeBERTa base model
 - ▶ bert-tiny
 - ▶ ALBERT base v2

Using a Task-Specific Model



Classification Tasks That Leverage Embeddings

- ▶ Typically a two step approach is used:
 - ▶ an embedding model for generating features.
 - ▶ a usual classifier fed with the features.



- ▶ Issue: generating embeddings of sentences

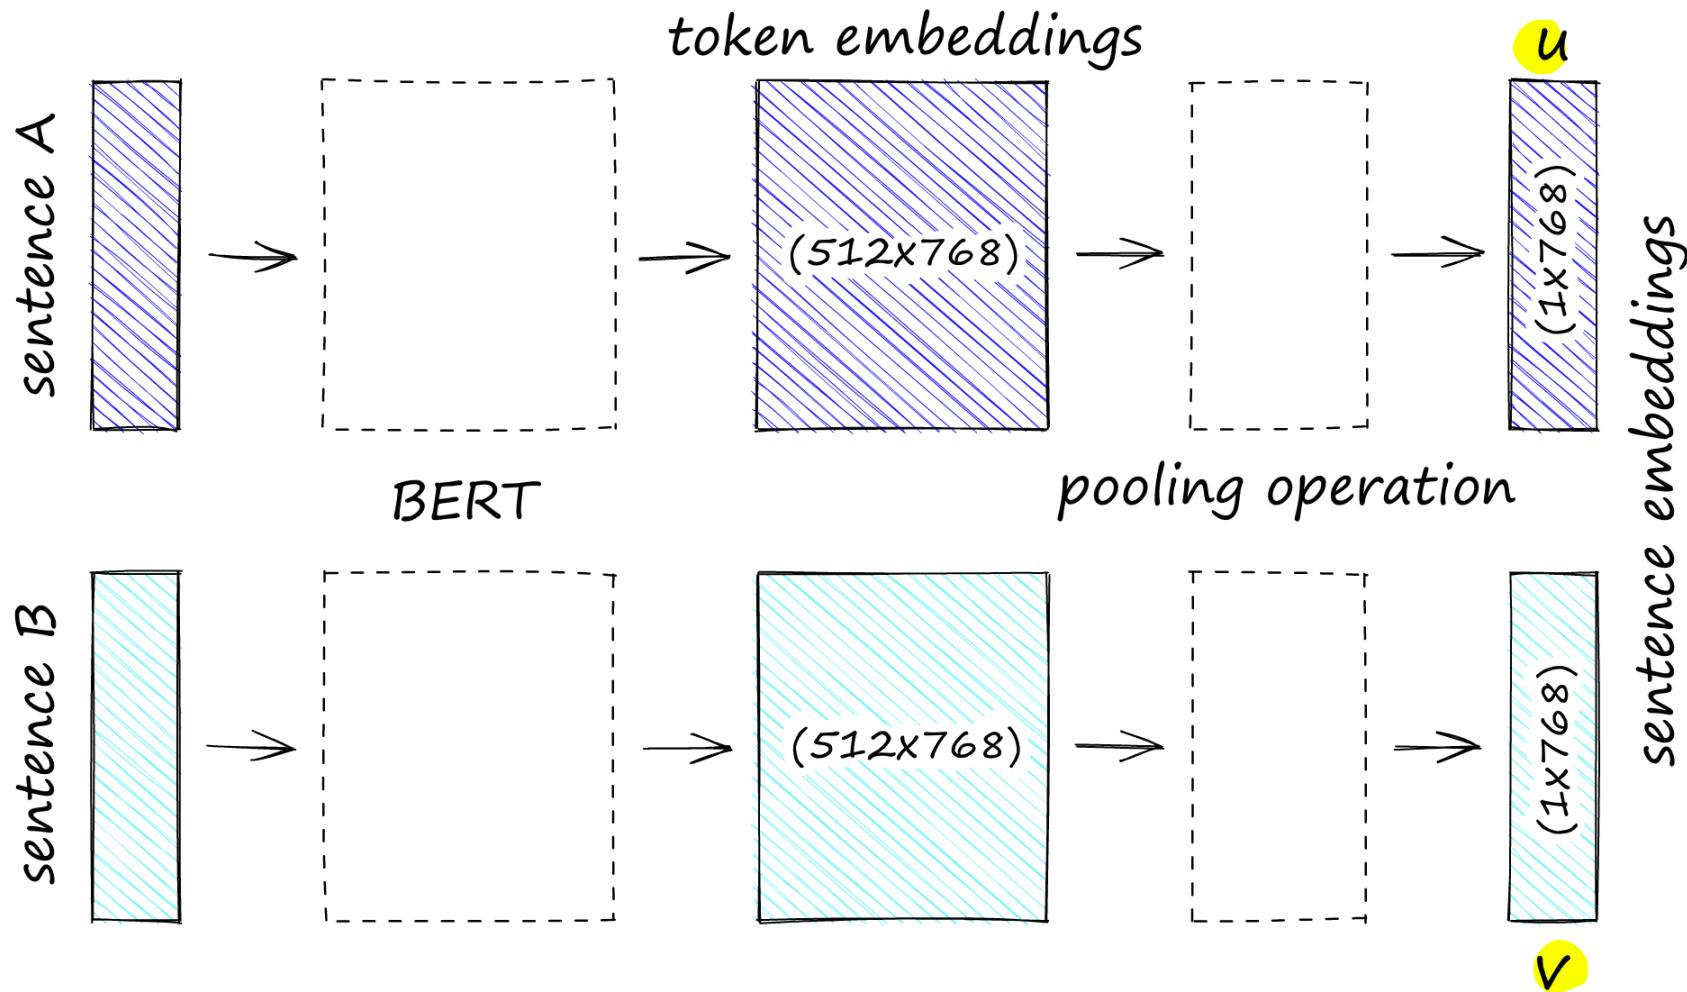
But... How Do We Generate Sentence Embeddings?

- ▶ The most intuitive approach:
 - ▶ Use BERT as an encoder to extract embeddings
- ▶ Two common strategies:
 - ▶ Use the [CLS] token
 - ▶ Average all token embeddings
- ▶ This gives us a **vector per sentence** usable in the 2-step pipeline.

Limits of BERT Embeddings

- ▶ BERT was not designed for independent sentence embeddings.
 - ▶ [CLS] and average pooling do not capture sentence-level semantics
 - ▶ They often perform worse than simple GloVe averages
- ▶ Embeddings are not positioned in vector space according to semantic similarity.
- ▶ ✓ Good for classification when trained jointly
- ▶ ✗ Poor when sentences are encoded independently.

Sentence-BERT: Sentence embeddings using Siamese BERT-networks.(EMNLP 2019)

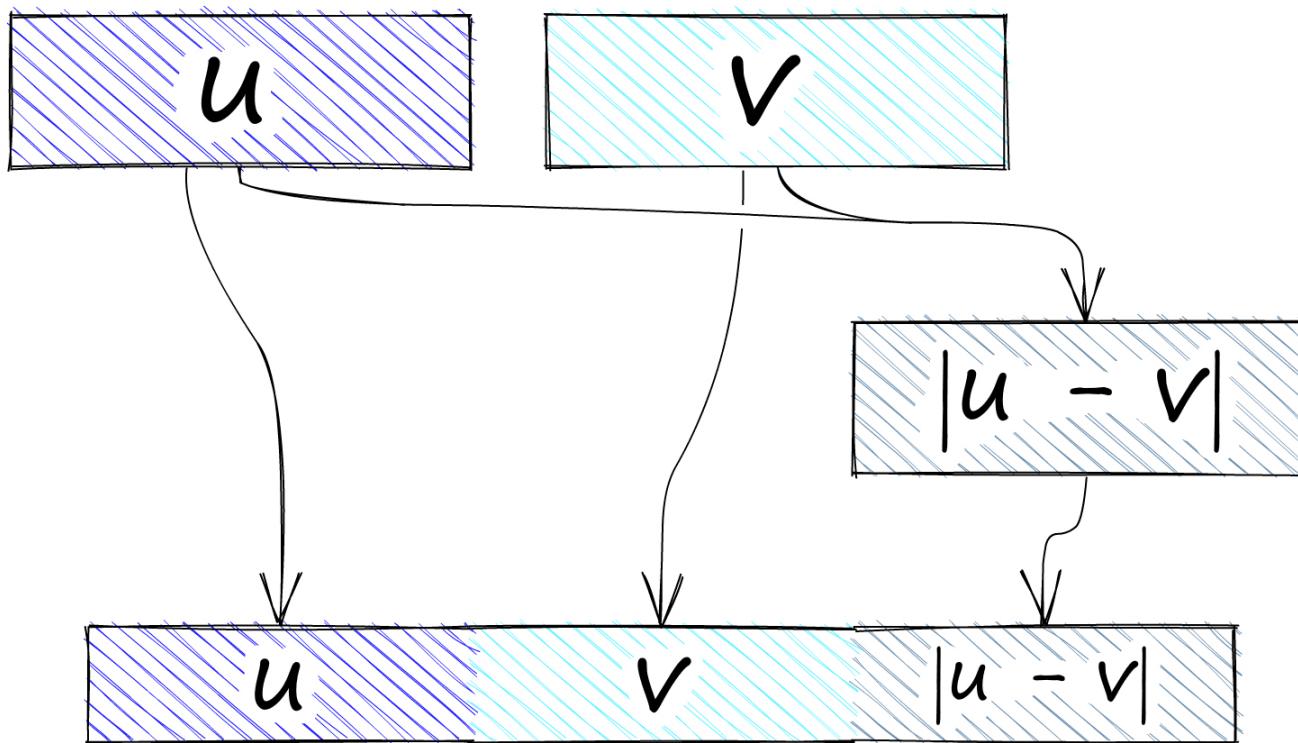


<https://www.pinecone.io/learn/series/nlp/sentence-embeddings/#Sentence-Transformers>

Sentence-BERT: Pre-training

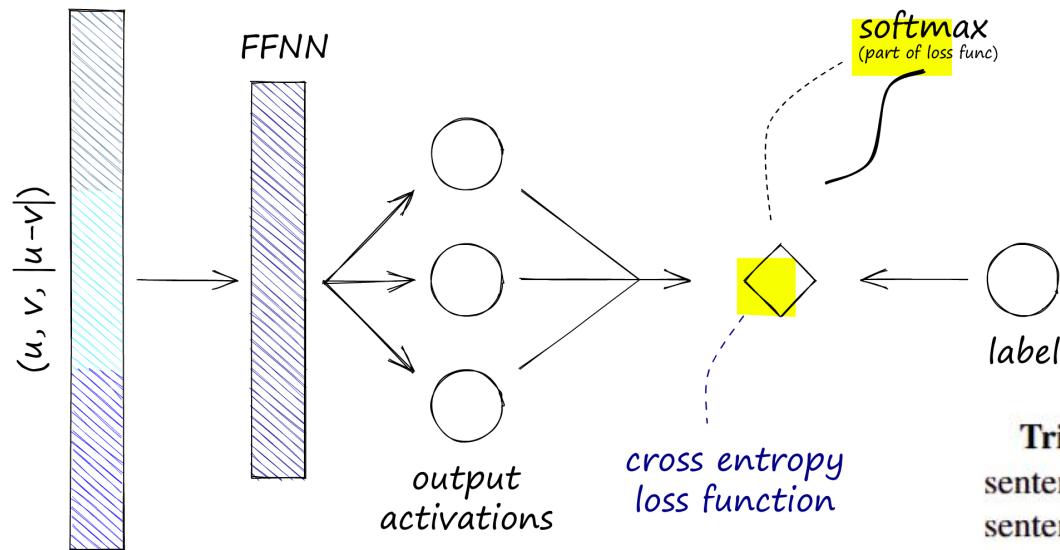
- ▶ **Objective:**
 - ▶ Fine-tune Siamese BERT using softmax-loss.
 - ▶ Datasets: SNLI (570K pairs) and MNLI (430K pairs).
- ▶ **Dataset labels:**
 - ▶ 0 – entailment: premise **implies** hypothesis.
 - ▶ 1 – neutral: premise and hypothesis **not necessarily related**.
 - ▶ 2 – contradiction: premise and hypothesis **contradict** each other.
- ▶ **Procedure:**
 - ▶ Feed sentence A (**premise**) into BERT A and sentence B (**hypothesis**) into BERT B.
 - ▶ Siamese BERT produces **pooled sentence embeddings**.
- ▶ **Pooling methods:**
 - ▶ **Mean pooling** (best performance for NLI and STS benchmark).
 - ▶ Max pooling
 - ▶ [CLS] pooling
- ▶ Francesco Guerra

Sentence-BERT: Pre-training



There are now two sentence embeddings: embeddings A u and embeddings B v . There are many concatenation approaches that we can test: highest performing is $(u, v, |u-v|)$.

Sentence-BERT: Pre-training



Classification Objective Function. We concatenate the sentence embeddings u and v with the element-wise difference $|u - v|$ and multiply it with the trainable weight $W_t \in \mathbb{R}^{3n \times k}$:

$$o = \text{softmax}(W_t(u, v, |u - v|))$$

where n is the dimension of the sentence embeddings and k the number of labels. We optimize cross-entropy loss. This structure is depicted in [Figure 1](#).

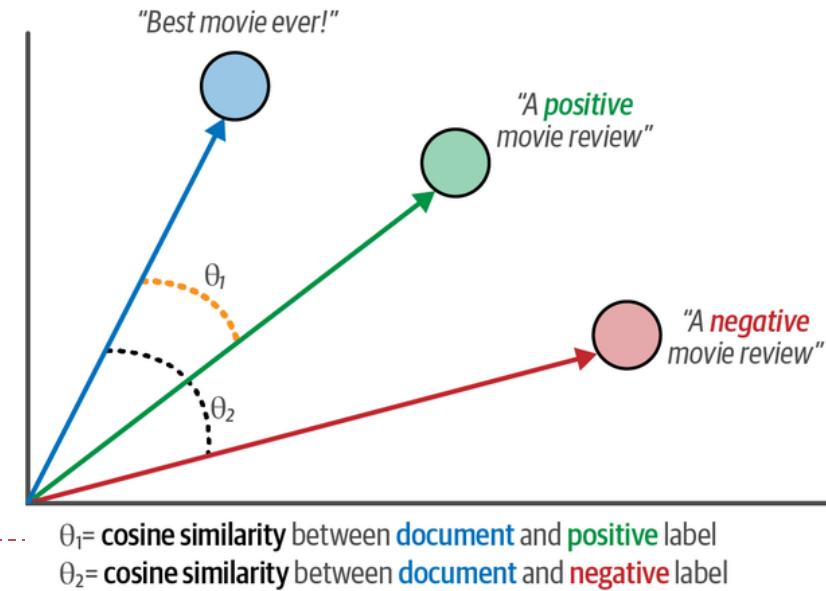
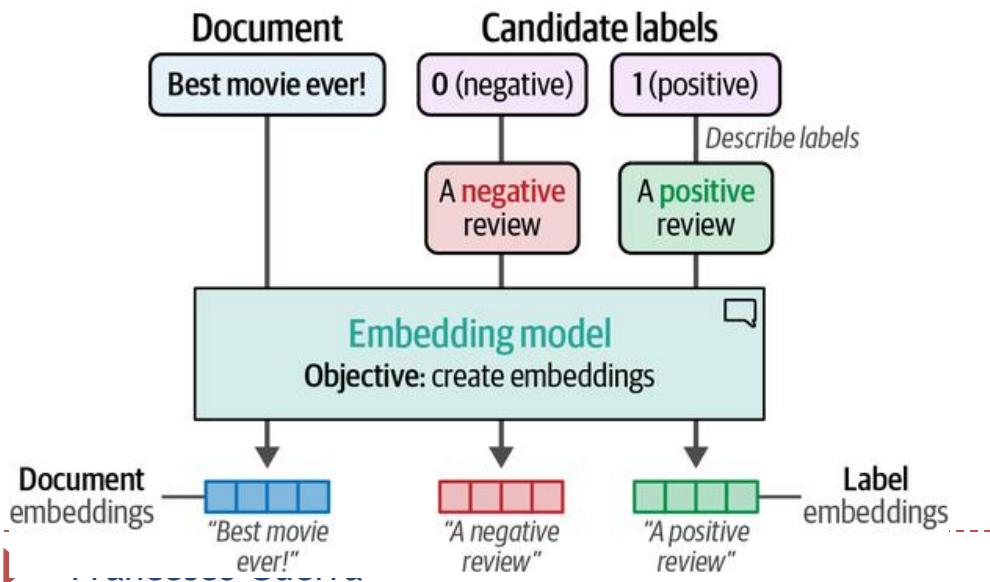
Triplet Objective Function. Given an anchor sentence a , a positive sentence p , and a negative sentence n , triplet loss tunes the network such that the distance between a and p is smaller than the distance between a and n . Mathematically, we minimize the following loss function:

$$\max(||s_a - s_p|| - ||s_a - s_n|| + \epsilon, 0)$$

with s_x the sentence embedding for $a/n/p$, $\|\cdot\|$ a distance metric and margin ϵ . Margin ϵ ensures that s_p is at least ϵ closer to s_a than s_n . As metric we use Euclidean distance and we set $\epsilon = 1$ in our experiments.

Zero-shot classification

- ▶ Zero-shot classification attempts to predict the labels of input text even though it was not trained on them
 - ▶ we can describe our labels based on what they should represent.
 - ▶ a negative label can be described as “This is a negative movie review.”
 - ▶ To assign labels to documents, we can apply cosine similarity to the document label pairs.



Text Classification with Generative Models

- ▶ Prompt engineering: iteratively improving your prompt to get your preferred output.
 - ▶ T5
 - ▶ GPT

Text Clustering



Text Clustering

- ▶ Text clustering aims to group similar texts based on their semantic content, meaning, and relationships, for
 - ▶ facilitating efficient categorization of large volumes of unstructured text
 - ▶ for quick exploratory data analysis.
- ▶ A possible pipeline:
 1. Convert the input documents to embeddings with an embedding model.
 2. Reduce the dimensionality of embeddings with a dimensionality reduction model.
 3. Find groups of semantically similar documents with a cluster model.

Dimensionality Reduction

- ▶ Transforming high-dimensional features into a lower-dimensional space preserving the most representative information.
 - ▶ PCA (Principal Component Analysis): Linearly projects data onto axes of maximum variance; assumes global linear structure.
 - ▶ T-SNE (t-Distributed Stochastic Neighbor Embedding): A nonlinear technique that preserves local similarities.
 - ▶ UMAP (Uniform Manifold Approximation and Projection): A learning technique that preserves both local and some global structures.
 - ▶ Autoencoder: Learns a nonlinear compressed representation through a neural network.

T-SNE key intuition

- ▶ Mapping high-dimensional data points into a two- or three-dimensional data, preserving the local relationships between points.
 - ▶ It achieves this by measuring the similarity between points in the high-dimensional space and representing this similarity as probabilities.
 - ▶ It constructs a similar probability distribution in the lower-dimensional space and minimizes the difference between distributions using gradient descent.
- ▶ Preserving local relationships between points
 - ▶ “Preserving the local relationships between points” refers to maintaining the relative distances and similarities between neighboring data points when they are mapped from a high-dimensional space to a lower-dimensional space.
 - ▶ Points that are **close in high-dimensional space** should remain **close in low-dimensional space**

Why Not Work Directly with Distances?

- ▶ Compute the distance between points is not reliable in higher dimensions due to the “curse of dimensionality.”
 - ▶ In high-dimensional spaces, **all distances become similar**
 - ▶ Difference between nearest and farthest neighbors shrinks
 - ▶ Distance rankings become unstable and noisy
- ▶ **Example Intuition**
 - ▶ In 2D: distances vary a lot → meaningful neighborhoods
 - ▶ In 100D: distances concentrate around the mean
 - ▶ Raw distances lose discriminative power

Step 1: Similarities in High Dimensions

- ▶ For each data point (e.g., x_1)
 - ▶ Compute distances to all other points
 - ▶ Convert distances into probabilities using a Gaussian distribution centered on the target point

$$P_{j|i} = \frac{\exp\left(-\frac{\|x_i - x_j\|^2}{2\sigma_i^2}\right)}{\sum_{k \neq i} \exp\left(-\frac{\|x_i - x_k\|^2}{2\sigma_i^2}\right)}$$

- ▶ Each point has its own σ_i
- ▶ This process is repeated for every point in the dataset, generating an $n \times n$ similarity matrix.
 - ▶ The resulting matrix is not symmetric, since the similarity from x_1 to x_2 may differ from x_2 to x_1 .
 - ▶ t-SNE symmetrizes them using $P_{ij} = \frac{P_{j|i} + P_{i|j}}{2N}$

Choosing Sigma (σ_i): Density Normalization via Entropy and Perplexity

- ▶ Data points may lie in regions with **very different densities**
- ▶ A fixed Gaussian width would:
 - ▶ oversmooth dense regions
 - ▶ undersmooth sparse regions
- ▶ Each point gets its own σ_i to adapt to its local density
- ▶ Role of Entropy
 - ▶ Quantifies how **spread out** the neighborhood probabilities are
 - ▶ Low entropy → few dominant neighbors (very local view)
 - ▶ High entropy → many neighbors with similar weight (broader view)
- ▶ Role of Perplexity
 - ▶ $Perplexity(P_i) = 2^{H(P_i)}$
 - ▶ Measures the **effective number of neighbors** that x_i considers close

How σ_i Is Computed (Binary Search)

- ▶ For each point x_i :
 - ▶ Start with an initial guess for σ_i
 - ▶ Compute conditional probabilities $P_{j|i}$
- ▶ Compute entropy $H(P_i)$ and perplexity
- ▶ Compare with the target (user-defined) number of neighbors to preserve
- ▶ Adjust σ_i :
 - ▶ increase $\sigma_i \rightarrow$ increase entropy / perplexity
 - ▶ decrease $\sigma_i \rightarrow$ decrease entropy / perplexity
- ▶ Repeat until perplexity matches the target

Example

- ▶ Number of neighbors to preserve = 3
- ▶ Initial guess $\rightarrow \sigma_i = 0.2$

Data point	Dist	Prob
A	0.10	0.543
B	0.20	0.374
C	0.40	0.083
D	0.80	~0
E	1.60	~0

- ▶ Entropy = 1.31 \rightarrow Perplexity = $2^{1.31} \approx 2.48$
- ▶ The perplexity suggests that less than 3 neighbors are taken into account...we have to adjust it
- ▶ $\sigma_i = 0.4 \rightarrow$ Perplexity = $2^{1.58} \approx 3$

Step 2: Similarities in Low Dimensions

- ▶ We now reduce the high-dimensional data into a lower-dimensional space, where data points are initially randomly placed along the x-axis.
- ▶ In this lower-dimensional space, we recalculate the similarity scores for each point relative to all others.
- ▶ This results in a second $n \times n$ similarity matrix, now based on low-dimensional distances.
- ▶ At this point, we have two matrices:
 - ▶ One representing similarity scores in the original high-dimensional space
 - ▶ One representing similarity scores in the lower-dimensional space

Optimization Objective

- ▶ The goal is to make the low-dimensional similarity matrix resemble the high-dimensional one as closely as possible.
- ▶ To do this, we minimize the difference between the two matrices using a divergence measure, typically the Kullback-Leibler (KL) divergence.
 - ▶
$$KL(P||Q) = \sum_{i,j} P_{ij} \log \frac{P_{ij}}{Q_{ij}}$$
- ▶ The algorithm iteratively adjusts the positions of points in the lower-dimensional space to reduce this divergence and better preserve the structure of the original data.
 - ▶ To minimize the KL divergence, we use gradient descent.

T-SNE configuration

- ▶ Perplexity: Measures the effective number of neighbors for each point. (5-50)
 - ▶ Balances local vs. global structure: Low perplexity → focus on local details High perplexity → captures global structure
- ▶ Learning Rate: Controls the step size during optimization. (10-100)
 - ▶ Too high → algorithm may oscillate or miss global minimum.
 - ▶ Too low → slow convergence, may get stuck in local minima.
- ▶ Number of Iterations: Defines how many times the algorithm updates the embedding. (1000)
 - ▶ Too few → embedding may not converge.
 - ▶ Too many → unnecessary computational cost.

UMAP in a nutshell

- ▶ Preserves Local & Global Structure: Captures both fine-grained details and overall data shape.
- ▶ Flexible Parameter Tuning: `n_neighbors` and `min_dist` control the locality-globality trade-off.
- ▶ Faster than t-SNE.
- ▶ High-level intuition
 - ▶ learns a graph in high-dimensional space
 - ▶ learns a similar graph in low-dimensional space
 - ▶ Goal: make the two graphs as similar as possible

UMAP Parameters

- ▶ **n_neighbors**
 - ▶ Controls the size of the local neighborhood used to build the graph structure (typical: 5-50)
- ▶ **min_dist**
 - ▶ Sets the minimum distance between embedded points.
 - ▶ Low (0.0-0.1): tight clusters
 - ▶ Medium (0.2-0.4): balanced
 - ▶ High (0.5+): spread-out structure
- ▶ **n_components**
 - ▶ Dimensionality of the target space (e.g., 2D or 3D).
- ▶ **Metric**
 - ▶ Distance function for computing neighbor similarity (e.g. euclidean, cosine, manhattan).

Step 1: Construct a k-Nearest Neighbor Graph

- ▶ For each point x_i :
 - ▶ find its k nearest neighbors (KNN)
- ▶ k is controlled by the parameter `n_neighbors`
- ▶ Interpretation:
 - ▶ Small $k \rightarrow$ focus on very local structure
 - ▶ Large $k \rightarrow$ incorporate more global structure

From Distances to Edge Weights

- ▶ For each point x_i and neighbor x_j :
 - ▶ Compute distance $d(x_i, x_j)$
 - ▶ Convert distances into weights $w_{ij} = \exp\left(-\frac{d(x_i, x_j) - \rho_i}{\sigma_i}\right)$
- ▶ Exponential function
 - ▶ Enforces rapid decay of similarity with increasing distance
 - ▶ Ensures that only very close neighbors have strong influence
- ▶ ρ_i = distance to the closest neighbor
 - ▶ zero distance to the nearest neighbor -> a **maximum weight of 1** for the nearest neighbor
 - ▶ Normalizes local neighborhoods across regions with different densities
- ▶ σ_i = local connectivity scale
 - ▶ Similar to t-SNE, σ_i adapts to local density
 - ▶ $\sum_{j \in KNN(i)} \exp\left(-\frac{d(x_i, x_j) - \rho_i}{\sigma_i}\right) \approx \log_2 k$

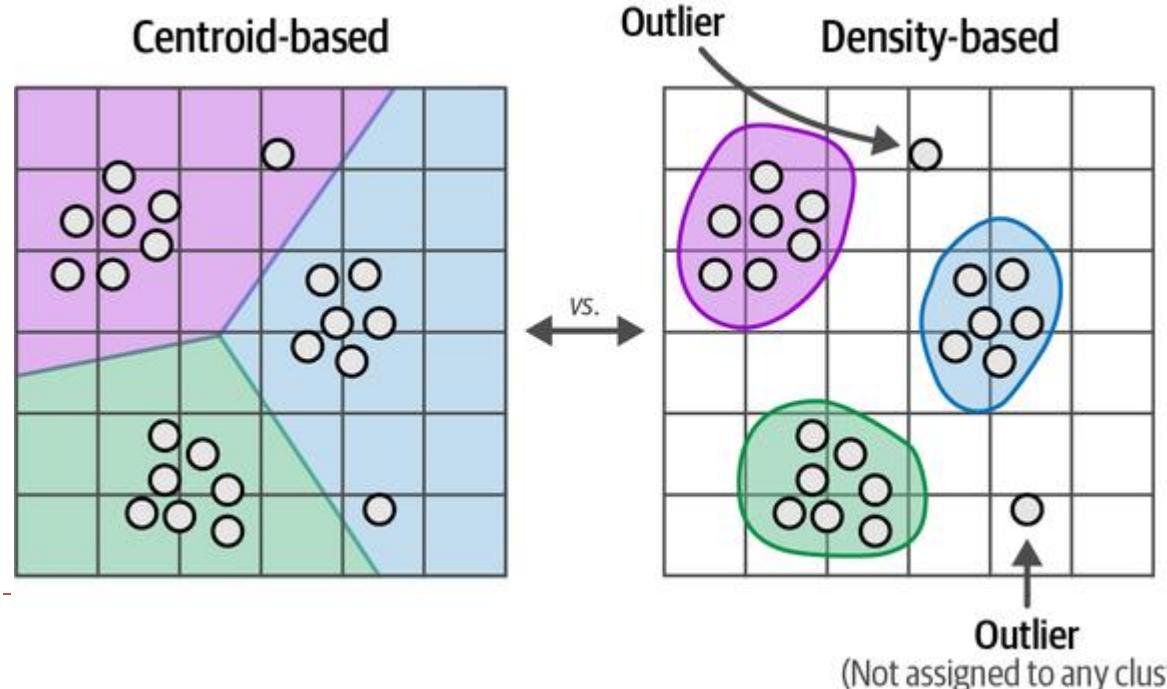
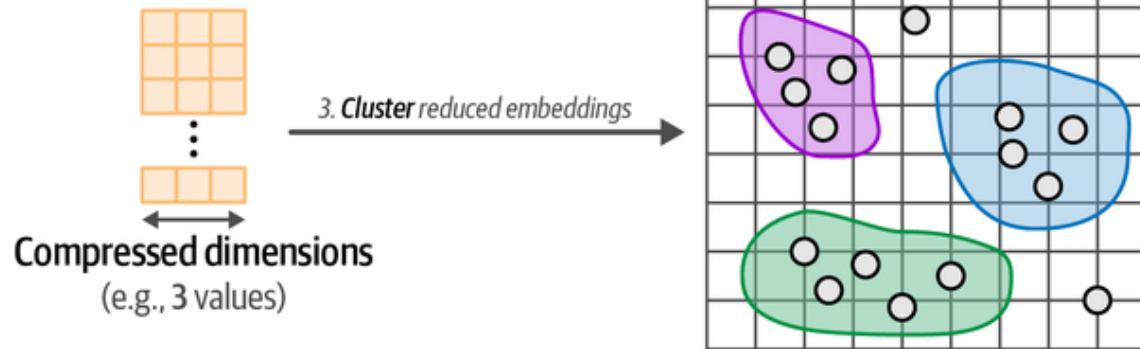
Symmetrization of the Graph

- ▶ Neighborhood relations are **asymmetric**
- ▶ UMAP symmetrizes edge weights using
 - ▶ $w_{ij}^{sym} = w_{ij} + w_{ji} - w_{ij} * w_{ji}$
- ▶ Preserves strong mutual connections
- ▶ Avoids shrinking weak one-sided links

Optimization Objective

- ▶ Minimize **cross-entropy** between high-D and low-D graphs
 - ▶ $CE = \sum_{i,j} w_{ij} \log(q_{ij}) + (1 - w_{ij}) \log(1 - q_{ij})$
- ▶ Encourages:
 - ▶ connected points to stay close
 - ▶ unconnected points to stay apart

Cluster the Reduced Embeddings



HDBSCAN in a nutshell

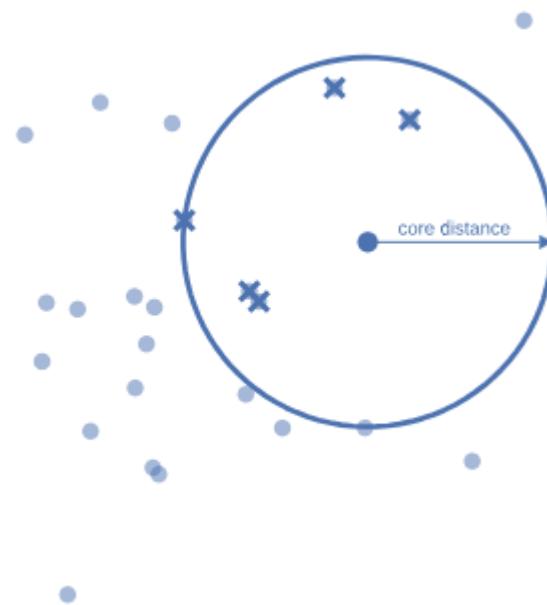
- ▶ HDBSCAN is a density-based clustering algorithm
 - ▶ Clusters are defined as regions with many nearby points
 - ▶ Sparse regions are considered separators or noise
- ▶ Intuition:
 - ▶ A cluster is a "crowded area"
 - ▶ Noise lives in "empty space"
- ▶ It automatically finds:
 - ▶ clusters of varying density
 - ▶ noise / outliers
- ▶ Widely used after UMAP or t-SNE embeddings
- ▶ No need to specify the number of clusters

HDBSCAN pipeline

1. Estimate local density (core distance)
2. Modify distances (mutual reachability)
3. Build MST
4. Construct hierarchy
5. Condense using minimum cluster size
6. Select clusters via stability

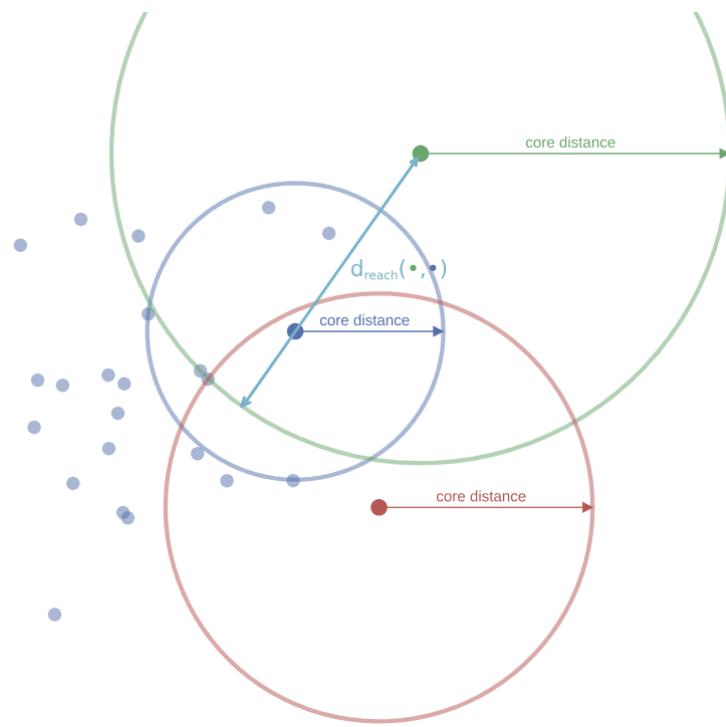
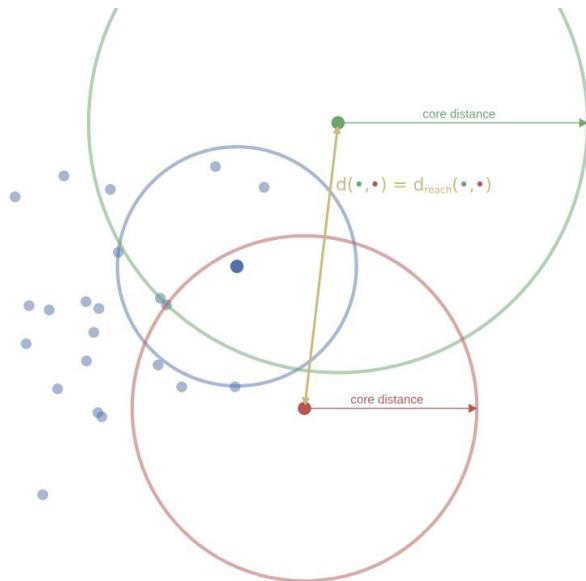
Step 1: Measuring Local Density

- ▶ For each point x_i :
 - ▶ find its k-th nearest neighbor (min_samples parameter)
- ▶ The distance to this neighbor defines the core distance:
- ▶ $\text{core_dist}(x_i) = \text{distance to the k-th nearest neighbor}$
- ▶ Dense regions → small core distance
- ▶ Sparse regions → large core distance



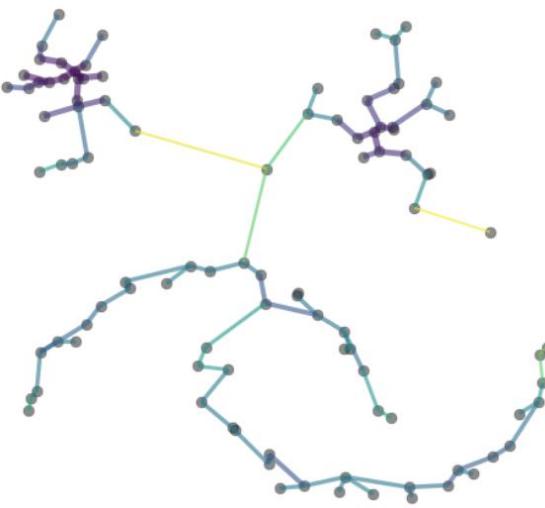
Mutual Reachability Distance

- ▶ To compare two points x_i and x_j , HDBSCAN uses:
 - ▶ $mreach(x_i, x_j) = \max(\text{core_dist}(x_i), \text{core_dist}(x_j), \text{distance}(x_i, x_j))$
- ▶ Intuition:
 - ▶ Penalizes pairs involving sparse points
 - ▶ Makes distances density-aware



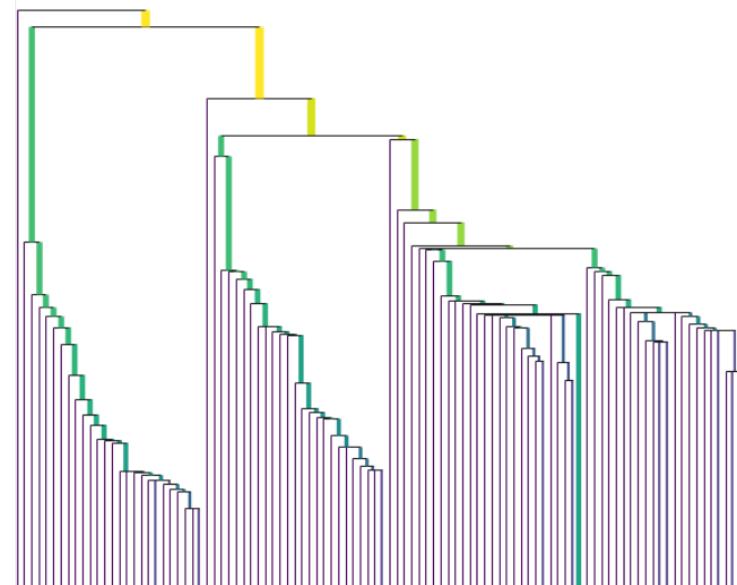
Build a Density-Aware Graph

- ▶ Construct a graph where:
 - ▶ nodes = data points
 - ▶ edge weight = mutual reachability distance
- ▶ Iteratively cutting a fully connected graph at different weight thresholds would be very expensive.
- ▶ Compute a Minimum Spanning Tree (MST) of this graph
- ▶ MST captures the essential density structure



From MST to Hierarchical Clustering

- ▶ Procedure:
 - ▶ Sort MST edges by increasing weight
 - ▶ Add edges one by one
 - ▶ Connected components merge over time
- ▶ This produces a **single-linkage dendrogram**:
 - ▶ Leaves = individual points
 - ▶ Merges happen at increasing distances



Why Distance-Based Cuts Are Not Enough

- ▶ Suppose we cut the dendrogram at a fixed distance:
 - ▶ Dense cluster → preserved
 - ▶ Sparse cluster → destroyed
- ▶ Or:
 - ▶ Sparse cluster → preserved
 - ▶ Dense cluster → merged with others
- ▶ **A single distance threshold cannot handle variable densities**

Reparameterization: From Distance to Density (λ)

- ▶ $\lambda = \frac{1}{distance}$
- ▶ Interpretation:
 - ▶ High $\lambda \rightarrow$ high density
 - ▶ Low $\lambda \rightarrow$ low density
- ▶ This does **not** change the tree structure, only how we *interpret* it.
- ▶ Now clusters live over **intervals of density**, not distance.

Cluster birth and death

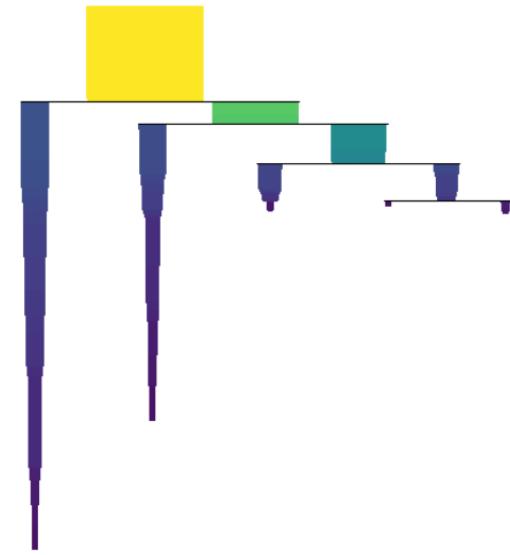
- ▶ A cluster is **born** when:
 - ▶ Its points first become connected
 - ▶ At a specific λ value
- ▶ A cluster **dies** when:
 - ▶ It merges with another cluster, or
 - ▶ It splits into valid subclusters
- ▶ Birth and death define the *lifetime* of a cluster.

Step 5: Condensed Cluster Tree

- ▶ Observation:
 - ▶ Many splits separate only 1-2 points
 - ▶ These are usually noise, not real clusters
- ▶ Introduce the parameter *minimum_cluster_size*
 - ▶ If a split creates a cluster smaller than this size → treat it as points falling out
 - ▶ Otherwise → accept it as a real split
- ▶ Result:
 - ▶ A simplified hierarchy: the **condensed tree**

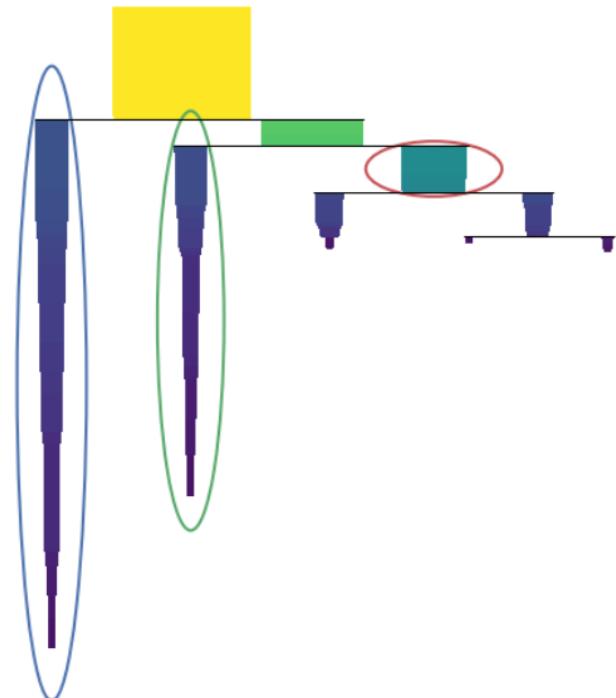
Example of condensed tree

- ▶ In the condensed tree:
 - ▶ Each cluster is drawn as a rectangle
 - ▶ x-axis: number of points (cluster size)
 - ▶ y-axis: density λ
- ▶ As λ decreases:
 - ▶ Points may fall out \rightarrow rectangle narrows
 - ▶ Cluster may split \rightarrow new rectangles appear



Cluster Stability

- ▶ Intuition:
 - ▶ A good cluster exists **for a long time** and contains **many points**.
- ▶ Stability measures:
 - ▶ How long a cluster persists
 - ▶ Weighted by how many points stay in it
- ▶ Formally:
 - ▶ $\text{stability} = \sum \text{over points of } (\lambda_p - \lambda_{\text{birth}})$
 - ▶ λ_{birth} : cluster birth
 - ▶ λ_p : point leaves the cluster
- ▶ Geometrically:
 - ▶ **Stability = area of the rectangle**



Step 6: Selecting the Final Clusters

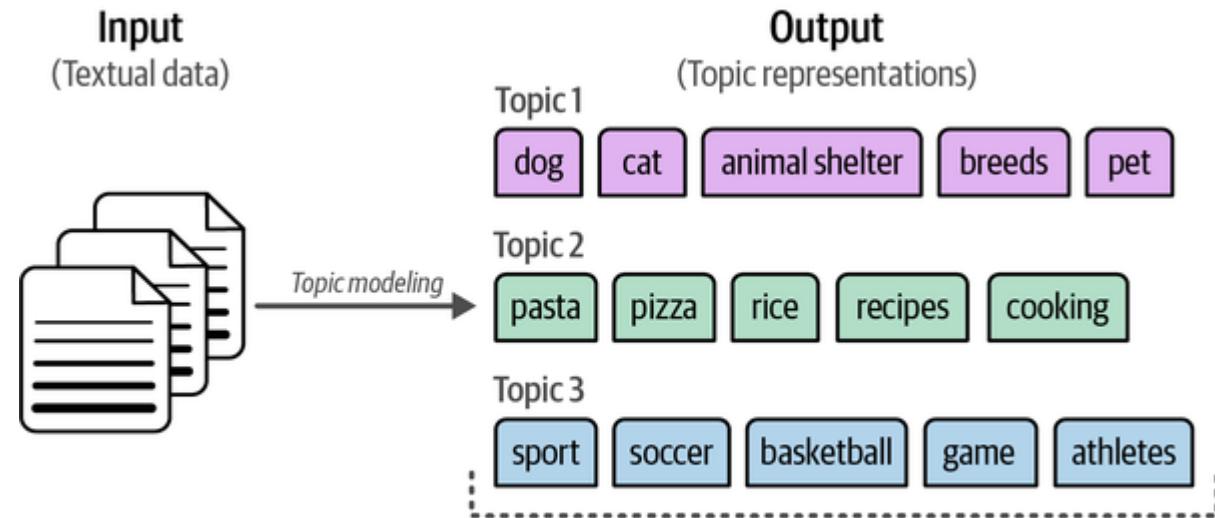
- ▶ We do not cut at a fixed λ .
- ▶ Instead:
 - ▶ Traverse the condensed tree bottom-up
 - ▶ For each cluster:
 - ▶ Compare its stability with the **sum of its children's stabilities**
- ▶ Rule:
 - ▶ If parent is more stable \rightarrow select parent
 - ▶ Otherwise \rightarrow select children
- ▶ This ensures:
 - ▶ No overlapping selections
 - ▶ Preference for robust structure

Final Output of HDBSCAN

- ▶ HDBSCAN returns:
 - ▶ **Cluster labels**
 - ▶ Points belonging to stable clusters get a label
 - ▶ **Noise points (label = -1)**
 - ▶ Points not consistently assigned become **noise**
 - ▶ **Membership strength** for each point
- ▶ Membership strength:
 - ▶ Derived from normalized λ values
 - ▶ Indicates confidence of assignment

Inspecting the Clusters / Topic modeling

- ▶ We can inspect each generated cluster manually and explore the assigned documents to get an understanding of its content.
- ▶ This idea of finding themes or latent topics in a collection of textual data is often referred to as *topic modeling*.
- ▶ Traditionally, it involves finding a set of keywords or phrases that best represent and capture the meaning of the topic

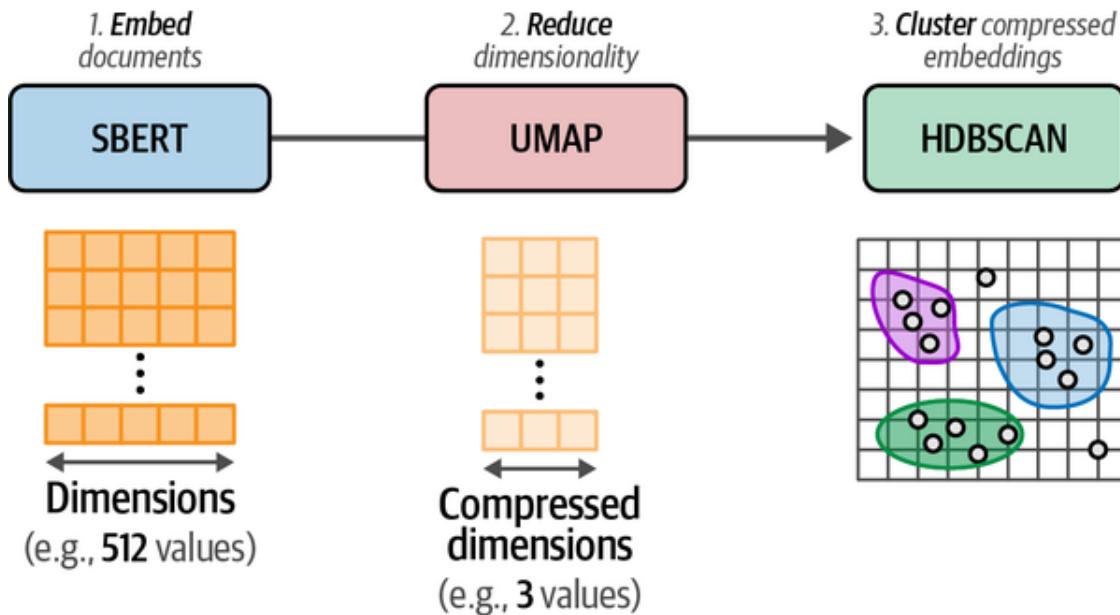


Topic Modeling

- ▶ Unsupervised statistical method for discovering hidden themes in large text collections (typically in the form of keywords).
- ▶ Automatically identifies topics from word patterns—no manual labeling needed.
- ▶ Reveals how themes are connected and how they evolve over time.
- ▶ Enables scalable organization and summarization of large text archives.

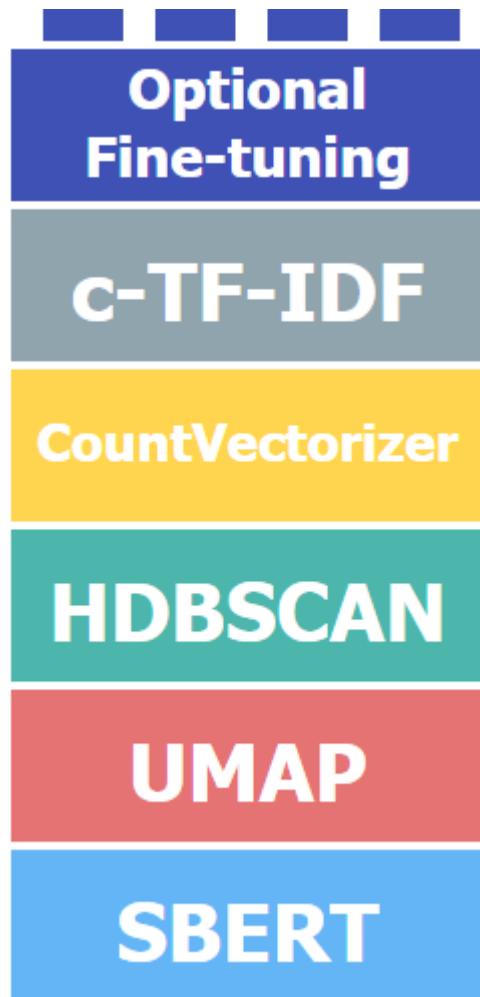
BERTopic in a nutshell

- ▶ Combines:
 - ▶ Transformer embeddings (semantic meaning)
 - ▶ Dimensionality reduction (UMAP)
 - ▶ Density-based clustering (HDBSCAN)
 - ▶ Interpretable topic representations (c-TF-IDF)



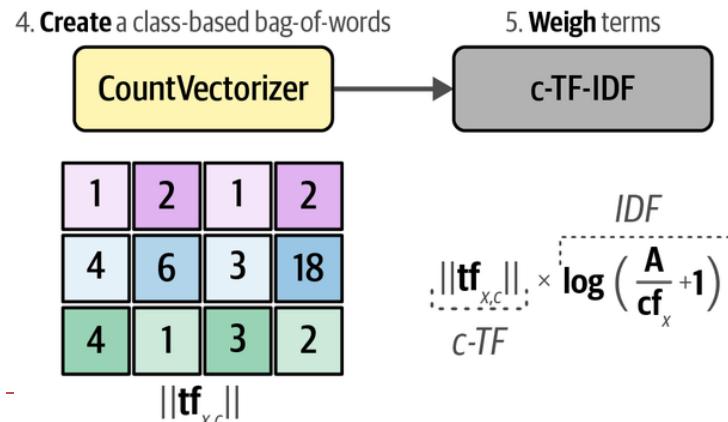
High-level pipeline

1. Embed documents with BERT / SBERT
2. Reduce dimensionality with UMAP
3. Cluster documents with HDBSCAN
4. Extract topic words with c-TF-IDF
5. (Optional) Refine topics with re-ranking



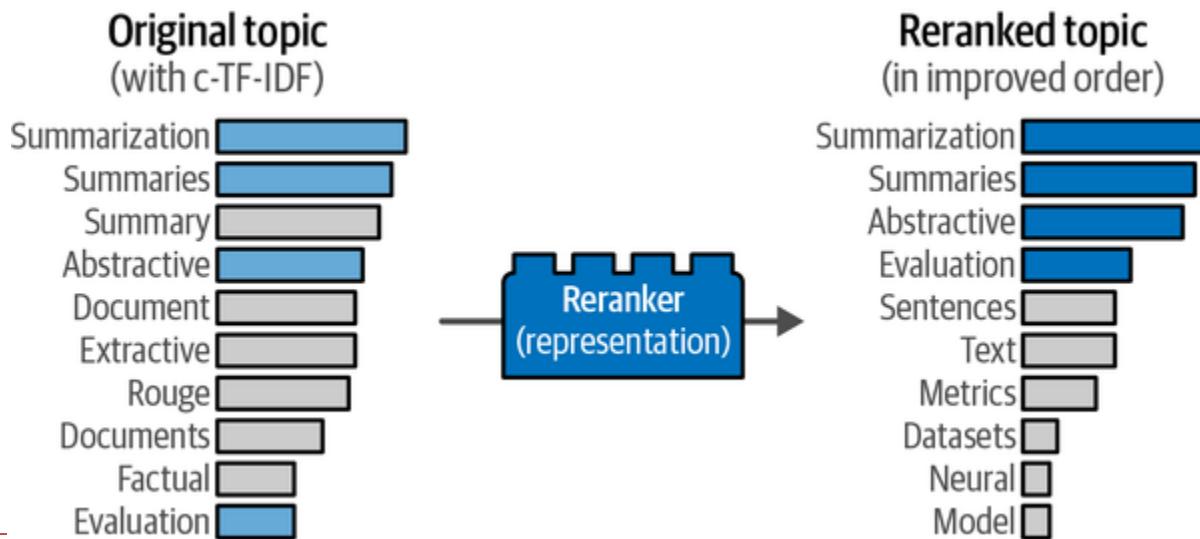
Topic Representation with c-TF-IDF

- ▶ Clustering gives groups of documents
- ▶ c-TF-IDF summarizes each cluster with keywords
 - ▶ a class-based variant of TF-IDF to put more weight on words that are more meaningful to a cluster and less weight on words that are used across all clusters.
 - ▶ $tf_{x,c}$: frequency of term x in cluster c (raw count)
 - ▶ cf_x : number of clusters where x appears
 - ▶ A : average number of words per cluster



Adding a Reranker

- ▶ The pipeline in BERTopic still represents a topic through a bag-of-words without taking into account semantic structures.
- ▶ A reranker leverages the strength of the bag-of-words representation to generate a meaningful representation.
 - ▶ We can use this first meaningful representation and tweak it using more powerful but slower techniques, like embedding models.





References

- ▶ D. Jurafsky & J. H. Martin: **Speech and Language Processing**, Pearson International 3ed draft <https://web.stanford.edu/~jurafsky/slp3/>
- ▶ D. Jurafsky, slides from NLP Courses <http://web.stanford.edu/~jurafsky/>
- ▶ Dan Klein, Slides from NLP Course
[http://www.cs.berkeley.edu/~klein/cs288/sp10/slides/SP10%20cs288%20lecture%203%20--%20language%20models%20II%20\(6PP\).pdf](http://www.cs.berkeley.edu/~klein/cs288/sp10/slides/SP10%20cs288%20lecture%203%20--%20language%20models%20II%20(6PP).pdf)
- ▶ J. Leskovec, A. Rajaraman, J. Ullman: Mining of Massive Datasets Stanford University <http://www.mmds.org>
- ▶ Landauer, T. K., Foltz, P. W., & Laham, D. (1998). Introduction to Latent Semantic Analysis. *Discourse Processes*, 25, 259-284.
- ▶ Bing Liu: Sentiment Analysis and Opinion Mining. *Synthesis Lectures on Human Language Technologies*, Morgan & Claypool Publishers 2012
- ▶ Jay Alammar, Maarten Grootendorst: Hands-On Large Language Models Published by O'Reilly Media, Inc.

References

- ▶ Zhijiang Guo, Michael Schlichtkrull, and Andreas Vlachos. 2022. [A Survey on Automated Fact-Checking](#). *Transactions of the Association for Computational Linguistics*, 10:178-206.